

Leonard Techel

Low-Connectivity State Space Exploration using Swarm Model Checking on the GPU

September 12, 2021

supervised by:

Prof. Dr. Sibylle Schupp
Sascha Lehmann

Hamburg University of Technology (TUHH)
Technische Universität Hamburg
Institute for Software Systems
21073 Hamburg

STS
Software
Technology
Systems

Abstract

Using small, independent verification tests, model checking large models with billions of states can be parallelized on GPUs. This approach works great on models with high connectivity. However, it fails when a model has only few edges between states or large portions of the state space are hidden behind bottleneck structures.

Past work on the *Grapple* model checker has tried different approaches including depth-limiting and alternating between breadth-first and depth-first search.

The main goal of this thesis is to: (1) create a systematic way of classifying a model as *low-connectivity* (2) minimize the amount of verification tests needed to maximize the state space coverage of said models. To do that, we provide an implementation of the *Grapple* model checker.

Contents

1	Introduction	1
2	Related Work	3
3	Background	5
3.1	Model Checking	5
3.1.1	Parallelized model checking	6
3.1.2	Swarm Verification	6
3.1.3	Grapple Framework	7
3.1.4	Low-Connectivity Models	7
3.2	CUDA	7
4	Theory	9
4.1	Model definition	9
4.2	Grapple model checker	9
4.2.1	Model Definition and State Generation	9
4.2.2	Queues	10
4.2.3	Hash table	10
4.3	State Space Diversification	12
4.3.1	Diversification Techniques	12
4.3.2	State Pruning and Start Overs	12
4.4	Waypoints model	12
4.5	Counting Unique States Visited	13
5	Implementation	15
5.1	Source code	15
5.1.1	Main program	15
5.1.2	CUDA kernel	15
5.1.3	Queues	15
5.1.4	Model and Successor Generation	16
5.1.5	Hash tables	17
5.1.6	HyperLogLog++	18
5.2	Usage	18
6	Experiments	21
6.1	Validating our implementation and reproducing results from the paper . .	21
6.2	Introducing Start Overs	22
6.3	Low-Connectivity Models	22
6.4	Optimizing Grapple on Low-Connectivity Models	23
7	Conclusion	25
7.1	Future work	25

7.2 Discussion	25
Bibliography	27

List of Algorithms

1	Fundamental State Space Exploration Loop	6
2	Grapple state space exploration loop of a single worker	11
3	Waypoints model	12
4	Branching-Free Ternary Operator	17
5	On-The-Fly State Generation on the GPU	18
6	Bitstate hashing	19

List of Figures

3.1	State pruning using hash collisions	7
3.2	Examples of low-connectivity properties	8
5.1	A single VT of the 4D array, mapped to 1D memory	16
6.1	Waypoints model, 1000 runs, 2^{18} hash table size, 0 start overs	21
6.2	Waypoints model, 1000 runs, 2^{14} hash table size, 15 start overs	22

1 Introduction

In an explicit-state model checker, state space exploration of large models with billions of states is a time-consuming problem.

Using swarm verification, the problem can be split into small, independent verification tests. Past work has shown that by executing these verification tests in parallel on GPUs, a high-speed model checker can be implemented [3].

GPUs are interesting for model checking because they allow a significant performance increase on parallel algorithms, are widely available and relatively cheap in comparison to specialized hardware like FPGAs/ASICs.

To create the small, independent verification tests, diversification is used. Each verification test only covers a subset of the state space. Together, the verification tests nearly achieve full state space coverage. This approach works great on models with high connectivity. However, it fails when a model has only few edges between states or large portions of the state space are hidden behind bottleneck structures.

2 Related Work

3 Background

This chapter gives a brief summary on the main topics that this thesis is based on. First, we define the scope of model checking that we are interested in. Then, we give an overview of Swarm Verification on the GPU and the specific problem that we are going to investigate: Low-Connectivity State Space Exploration. Last but not least, we give an introduction to the CUDA GPU programming framework.

3.1 Model Checking

Model checking is a formal verification method that checks whether a state machine satisfies a specification. For example, the state machine of an elevator may be verified to meet the safety property of not opening the doors between floors. To do so, a *model checker* searches the state space for counterexamples, also called violations. When a violation is found, the path of state transitions that lead to the violation is reported back.

Model checking can be divided into three main problems: Description of models through state machines, definition of the specification through temporal logic, and algorithms that verify whether a state machine models a specification.

There are two main branches of model checking: Explicit-State Model Checking and Symbolic Model Checking. Explicit-State Model Checking can only verify finite state machines. In particular, the model has to have finite states, each state needs to be representable by a finite-size tuple containing its atomic propositions and the model changes state through execution of state transitions. To overcome these limitations and verify potentially infinite-size state machines or systems of unknown structure, Symbolic Model Checking uses the abstraction of *symbols*, each representing a set of states and transitions. Within this thesis, we are only considering explicit-state model checking.

Each state in a model is labelled with atomic propositions that hold true while the state is active. An example for such propositions are the current values of variables in a program at a given state. In a specification, different types of properties can then be expressed onto these propositions. Three common properties are reachability, safety and liveness: Reachability means that an atomic proposition holds true at some state in the future. Safety means that an atomic proposition holds true at all states in the future. Liveness means that an atomic proposition holds true infinitely often in the future, meaning that it does not happen that the atomic proposition never holds true. Within this thesis, we are only considering reachability and safety properties.

A challenge all model checking algorithms have to face is the *state explosion problem*. In an asynchronous model of n processes, each consisting of m states, the number of states grows exponentially by the number of processes, namely m^n . This means that even for small models, it is often not possible to fit all reachable states of the system into a computer's memory. Therefore, every model checking algorithm needs to reduce the state space in some sense. However, even then, a non-parallel algorithm may need a lot of physical time for exhaustive verification of the state space. In exhaustive verification, all states are visited and checked for a violation. [2, 6]

3.1.1 Parallelized model checking

The basic operation in an explicit-state model checker is a state space exploration loop, as defined in Algorithm 1. Model checking can be speed up by parallelizing the state space exploration, for example using the parallel breadth-first search (BFS) from the paper “Parallelizing the Spin Model Checker” by Holzmann [7].

A major challenge in parallelized BFS is the communication overhead between threads: Shared memory does not allow to easily split the work onto a cluster of heterogeneous processors or the massively parallel architecture of a GPU on which thousands of threads can exist simultaneously.

Algorithm 1 Fundamental State Space Exploration Loop

```

while there are unvisited states do
  mark state as visited
  if state violates spec then
    report path to state

```

3.1.2 Swarm Verification

A new approach on parallelized model checking comes from the paper “Swarm Verification” by Holzmann et al. [8]. Swarm Verification solves the challenge of parallelizing state-space search by splitting the state space exploration into many small, independent, memory-limited tasks called *Verification Tests* (VTs). Each VT only covers a small subset of the total state space and, using *diversification techniques*, uses a different search path. By executing all VTs and collecting their results, we still achieve nearly full state space coverage.

As VTs are independent of each other, we can easily execute them on heterogeneous computers. Even further, as VTs are also memory-limited, we can massively parallelize them on devices with very limited resources like GPUs.

Multiple diversification techniques can be applied, e.g. randomizing the order of nondeterministic choice on states with multiple outgoing transitions or reversing the search order. The most powerful diversification technique is state pruning using hash collisions: During state exploration, states get marked as visited in a limited-size hash table. By using a different hash function for each VT, the state space gets automatically pruned through hash collisions. The process is illustrated in Figure 3.1, that shows an example state graph on which BFS is performed by two VTs. On the left side, state B and C cause a hash collision, resulting in the subgraph originating in state C being removed. On the right side, state E and F cause a hash collision, resulting in the subgraph originating in state F being removed. Each state exploration on its own does not cover the whole state space, however, by combining the results from both searches, all states are covered again.

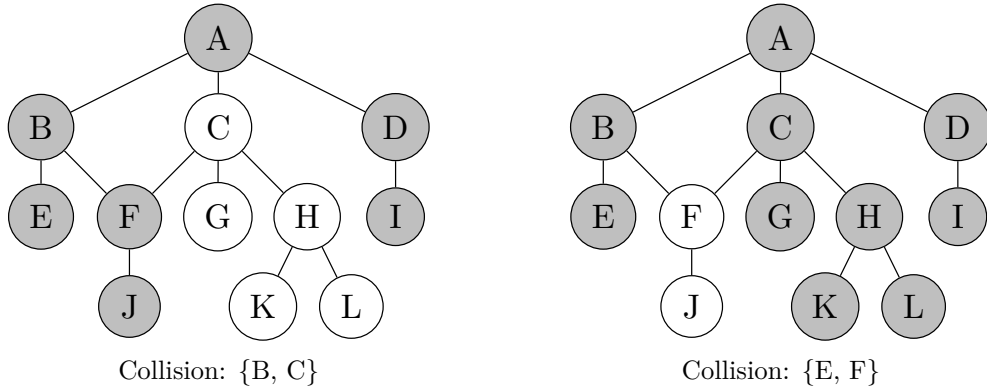


Figure 3.1: State pruning using hash collisions

3.1.3 Grapple Framework

The paper “Swarm model checking on the GPU” by DeFrancisco et al. [3], on which this thesis is based, contributes the *Grapple* framework for parallel swarm verification on the GPU using CUDA. In their implementation of Swarm Verification, each VT’s state space exploration runs in a parallel BFS, using CUDA’s built-in synchronization primitives to coordinate threads.

3.1.4 Low-Connectivity Models

A key observation of [3] is that their algorithm’s state space coverage in relation to the number of executed VTs slows down on models with *low connectivity*. A model has *low connectivity* if at least one of the following properties is satisfied:

- **Generally Linear:** The average number of edges per state is close to two: One inbound, one outbound
- **Bottleneck Structures:** A single state or group of states other than the initial state that needs to be passed to reach most of the state space

See Figure 3.2 for examples of these properties.

3.2 CUDA

CUDA is NVIDIA’s proprietary GPU programming framework that provides automatic massive scalability.

The CUDA model defines a three-level abstraction. At the lowest level, a program function called *kernel* is defined using C/C++. The kernel is executed by defining the amount of *threads* and *blocks* that run in a *grid*: Each grid contains multiple blocks. Each block contains multiple threads. Each thread in a block executes the kernel function using the single instruction, multiple thread (SIMT) execution model. Thus, each thread should




Figure 3.2: Examples of low-connectivity properties

work on an independent piece of memory. Only threads within a block can cooperate, for example using shared memory or synchronization barriers.

Internally, the GPU, also called *device*, maps the blocks of threads onto its *streaming multiprocessors*. By doing so, tens of thousands of threads can be executed in parallel. The architecture allows scaling even further onto multiple GPUs.

Furthermore, CUDA defines multiple memory levels, each with different size and access speed.

Registers are used for thread-local variables. There can be at max 255 registers à 32 bit per thread.

Shared memory can only be accessed by the threads of the block in which it got allocated. As it is on-chip, it is faster than global memory. By default, there is 48 KB of shared memory.

Constant memory is part of the device memory, with the constant cache in front of it. Cache hits are served at higher speed than just using global memory. There is 64 KB of constant memory.

Global memory is the largest memory with usually multiple gigabytes of size. It is placed in device memory and thus also the slowest.

CUDA requires the problem to be solved by a program to be split into many small sub-problems. This is a good fit with Swarm Verification, which turns the model checking problem into many small, independent VTs.

4 Theory

This chapter gives a detailed explanation of the algorithms used by our model checker. It starts with an in-depth description of the Grapple framework’s state space exploration loop and its components in Section 4.2. We then continue with a list of known State Space Diversification Techniques in Section 4.3. The chapter finishes with a solution to the problem of counting the number of unique states visited in Swarm Verification in Section 4.5.

4.1 Model definition

The model is defined using a *Kripke Structure* $M = (S; S_0; R; L)$ where S is the finite set of states, $S_0 \subseteq S$ is the finite set of initial states, $R \subseteq S \times S$ is the left-total transition relation and $L : S \mapsto 2^{AP}$ is the labelling function that assigns each state its atomic propositions.

4.2 Grapple model checker

In Grapple, each VT executes an internally parallel state space exploration. We call it *internally parallel* to differentiate between the parallelism *within* a VT and the parallelism *between* multiple VTs.

Each thread in a VT executes Algorithm 2, which is a more abstract version of Algorithm 1 from “Swarm model checking on the GPU” by DeFrancisco et al. [3]. We are going to call this algorithm the *Grapple algorithm*. The Grapple algorithm is based on the parallel BFS from “Parallelizing the Spin Model Checker” by Holzmann [7].

4.2.1 Model Definition and State Generation

State generation strategies of model checkers can be divided into *a priori* and *on-the-fly* generation. In a priori generation, the state space graph is completely known before running the verification algorithm, for example as an adjacency matrix. In on-the-fly generation, the successors of each state are created on the fly during verification. Grapple uses on-the-fly generation on the GPU. Furthermore, Grapple can only verify models with the same number of nondeterministic choices

$$M = ((S; S_0; R; L); P; NDC)$$

The states of a model support two operations: `state_successor($p, ndc, state$)` on-the-fly returns a successor to a state and `state_violates($state$)` returns whether a state violates.

$$\text{state_successor} : S \times NDC \times P \mapsto S$$

$$\text{state_violates} : S \mapsto \{0, 1\}$$

4.2.2 Queues

In summary, the parallel BFS is implemented using lock-free communication between threads via a shared, multidimensional queue array. Being lock-free removes the overhead of waiting for locks to be released and allows the algorithm to be split up onto many threads, limited only by the quadratic queue size.

For each pair of threads, there is only one piece of memory to communicate over. Writing and reading to this piece of memory is coordinated by splitting the algorithm into two alternating phases:

In the first phase $t = 0$, each thread i processes all states from its input queues $Q[t][0 \dots N][i]$. Here, N denotes the amount of threads executing a VT, with $N = 32$ being the default. For each input state, successors are generated on-the-fly. For each successor, we check whether it is already visited in a hash table. If a successor is priorly unvisited, we check whether it violates. Violating successors are reported to the host using a buffer. Non-violating successors are written to a random worker j 's input queue $Q[1 - t][i][j]$. When all input states are processed and threads are synchronized, the phases get swapped, i.e. $t = 1 - t$. This process repeats until no more unvisited successors are discovered.

The queues support three operations: `queue_push($Q, state$)` adds a state to the back, `queue_pop(Q)` removes and returns the first state from the front and `queue_empty(Q)` tells whether a queue is empty.

4.2.3 Hash table

The hash tables support one operation: `mark_visited($V, state$)` marks a state as visited and returns whether it was already visited before.

Completeness A big challenge in Swarm Verification is the question on how much of the state space is actually covered, i.e. whether all violations are found. Due to overlap in state space between VTs and the sheer amount of states being checked by each VT, computing the exact number of unique states visited and thus the exact state space coverage can hardly be done. Instead, we can estimate the state space coverage, as described in Section 4.5.

Termination The algorithm terminates when there are no more unvisited state successors according to the hash table. In this case, all input queues of the next phase remain empty, i.e. $Q[1 - t][0 \dots N][0 \dots N] = \emptyset$.

There are two reasons why at some point no more unvisited states get discovered: Either, we did an exhaustive verification of the model, meaning that all states actually got visited. Or, which happens most of the time in Grapple, the hash table is full, causing

a collision of every new successor with an already visited state. Collisions are intended, as mentioned in Section 3.1.2.

Resulting from this, the algorithm at worst terminates after having visited the number of states being the minimum of queue size and state space size.

Algorithm 2 Grapple state space exploration loop of a single worker

```

 $t \in \{0, 1\} \leftarrow 0$   $\triangleright$  Current algorithm phase
 $visited \subseteq S \leftarrow \emptyset$   $\triangleright$  Hash table of visited states
 $Q[2][N][N][I] \leftarrow \emptyset$   $\triangleright$  Queues

 $\triangleright$  Initial state  $\triangleleft$ 
queue_push( $Q[t][0][0], S_0[0]$ )
mark_visited( $S_0[0]$ )

__syncthreads()

done  $\leftarrow$  false
while not done do
  for  $i = 0, \dots, N$  do
    while not queue_empty( $Q[t][threadIdx.x][i]$ ) do
      state  $\leftarrow$  queue_pop( $Q[t][threadIdx.x][i]$ )
      for  $p \in$  processes do
        for  $ndc \in$  nondeterministic choices within  $p$  do
          succ  $\leftarrow$  state_successor( $p, ndc, state$ )
          visited  $\leftarrow$  mark_visited(succ)
          if not visited then
            if state_violates(succ) then
              report path to state
            else
              next  $\leftarrow$  random output queue  $n \in N$ 
              queue_push( $Q[1 - t][next][threadIdx.x], succ$ )

      __syncthreads()

  done  $\leftarrow$  queue_empty( $Q[1 - t][0 \dots N][threadIdx.x]$ )
   $t \leftarrow 1 - t$ 

  __syncthreads()
  
```

4.3 State Space Diversification

4.3.1 Diversification Techniques

These diversification techniques exist:

1. State Pruning using Hash Collisions
Each VT marks its visited states in a limited-size hash table with varying hash functions. Hash collisions cause different partitions of the state space. Hash table exhaustion causes termination of the exploration when all new states are supposedly already visited.
2. Reversing search direction or order
3. Randomizing order of nondeterministic choice
4. Parallel Deep Search (PDS)
5. process-PDS

4.3.2 State Pruning and Start Overs

4.4 Waypoints model

The Waypoints (WP) model is a benchmark for Swarm Verification model checkers. It is first introduced in [9] and is used as primary benchmark in [3]. Our implementation is defined in Algorithm 3.

The model creates a state space of the size of all 32-bit integers, i.e. $2^{32} = 4\,294\,967\,296$ states. To do so, it uses eight processes, each in control of four bits. At successor generation, each process will nondeterministically set one of its bits.

There is a single reachability property: We predefine a set of 100 uniform randomly chosen 32-bit integers, called *waypoints*. A state violates if it is part of the set. As the waypoints are uniformly distributed, the number of unique discovered waypoints is equal to the percentage of state space coverage. For example, 24 discovered waypoints means that approximately 24% of the state space is covered.

Algorithm 3 Waypoints model

$violations \leftarrow \{100 \text{ random } 32\text{-bit integers}\}$

function state_successor($p, ndc, state$)
 | **return** $state \mid 1 \ll ((4 \cdot process) + ndc)$

function state_violates($state$)
 | **return** $state \in violations$

4.5 Counting Unique States Visited

In Swarm Verification, we achieve a much faster verification by only covering *nearly* 100% of the state space. Resulting from this, it is important to know how much of the state space is actually covered. To find this out, we need to count unique states visited.

To calculate the percentage of state space covered, we have to know the total state space size. As in this thesis we only consider models with known state space size, we can provide it as pre-calculated constant.

The executed VTs may overlap in explored state space, meaning that across all VTs, a state may be visited multiple times. In order to calculate the exact number of unique states visited, we have to identify distinct states in the stream of all visited states. This is called the *count-distinct problem*.

Intuitively, one could collect all visited states in a set and then take its cardinality. However, this approach requires an amount of memory proportional to the amount of visited states which, as of the state space explosion problem, results in exponential memory usage.

A solution to this problem are *probabilistic cardinality estimators* that approximate the number of unique states within a fixed error using significantly less memory.

We choose the algorithm introduced in “HyperLogLog: the analysis of a near-optimal cardinality estimation algorithm” by Flajolet et al. [4], as it is commonly used and relatively easy to implement.

For each VT, we create a HyperLogLog. Inside each VTs state space exploration, we call the **add** operation for each newly discovered state. When a VT has finished, we **merge** its HyperLogLog with those of all other already finished VTs. We then can execute the **estimate** operation on the global HyperLogLog to get the estimation of unique visited states across all finished VTs.

Cardinality estimation with HyperLogLog yields good results on the Dining Philosophers problem with 15 processes and $3^{15} - 1$ states. As expected, the number of unique states visited grows logarithmic until reaching around 100%. However, on the Waypoints model with 2^{32} states, the estimation fails.

The paper “HyperLogLog in Practice: Algorithmic Engineering of a State of the Art Cardinality Estimation Algorithm” by Heule et al. [5] finds out that the estimation of cardinalities beyond one billion fails on the original HyperLogLog algorithm. They solve this problem by using a 64-bit hash instead, calling their improved algorithm HLL++.

Using HLL++, we can count unique states visited.

5 Implementation

Our implementation serves two purposes: First, to reproduce the results from [3]. Second, as foundation for experimenting with the state space exploration of low-connectivity models.

This chapter documents our implementation of a Grapple model checker. The general architecture of the model checker is described in Section 5.1. Usage instructions are described in Section 5.2.

5.1 Source code

We chose C++17 as programming language, so we can make full use of the CUDA Toolkit, which is provided as C header. The CUDA Toolkit is used in version 11.4. As build system, CMake 3.16 is used. Code is written using the Object-Oriented Programming paradigm.

Our implementation consists of six components: 1. The main program, running the CUDA kernels and collecting their results. 2. The CUDA kernel, implementing a parallel state space exploration loop. 3. The queues, providing lock-free communication between threads in a VT. 4. The hash tables in which states are marked as visited. 5. The model, providing successor generation and violation checking. 6. The HyperLogLog, counting unique states visited across all VTs.

In the following, we are going to describe the implementation-specific details and challenges of each component.

5.1.1 Main program

The main program runs in a single thread on the host. On startup, it seeds a global pseudorandom number generator (PRNG) and creates a random value for each CUDA thread. Then, for each run, it executes the CUDA kernel in a grid of $K = 250$ blocks, each consisting of $N = 32$ threads. Each block represents a VT, meaning that VTs are executed in batches of 250. Each thread represents a worker of a VT's parallel state space exploration. After each run, the main program collects the discovered violations and number of unique states visited, accumulates them and prints them to the standard output line-by-line as CSV.

5.1.2 CUDA kernel

The CUDA kernel executes the state space exploration loop, as described in Section 4.2.

5.1.3 Queues

The queues provide a data structure for lock-free communication between threads in a VT, as described in Section 4.2.2. By default, each queue has a capacity of $I = 4$ states. They are only used on-device within the CUDA kernel and stored in global

memory due to the limited size of shared memory. Resulting from this, we have to map $K \times 2 \times N \times N \times I$ queues into memory. This proposes two challenges: 1. We have to implement fixed-capacity queues. 2. We have to map multidimensional queues into a one-dimensional, linear memory allocation.

Usually, fixed-capacity queues are implemented as ring buffer. However, due to the design of the two-phase parallel state space exploration, it is sufficient to implement the queue as a singly linked list. In addition to that, we have to keep a pointer to the **head** and **tail** of the linked list. Then, in the first phase, each queue is filled through the `queue_push` operation by adding an element to its linked list and updating the **tail**, until the **tail** points to the last element in memory. In the second phase, each queue is completely emptied through the `queue_pop` operation by retrieving the **head**, then incrementing it until **head** equals **tail**. A queue is empty if both **head** and **tail** are null pointers.

The memory address of thread i 's input queue, which is an output queue of thread j , in algorithm phase t , of VT v , is calculated using the formula:

$$v \cdot (2 \cdot N \cdot N) + t \cdot (N \cdot N) + j \cdot N + i$$

Figure 5.1 illustrates the mapping of a single example VT with $K = 1$ and $N = 4$, meaning that we allocate memory for $1 \times 2 \times 4 \times 4 = 32$ queues. On the top row, each slot represents the memory allocation of a single queue. The remaining rows illustrate the constants enclosed in parentheses in our formula: Each set of output queues has a width of N , which we have to multiply with the index of the current thread to get at its first memory slot. Each phase has a width of $(N \cdot N)$, which we have to multiply with the index of the current phase. Each VT has a width of $(2 \cdot N \cdot N)$, which we again have to multiply with the index of the current VT to get to its first memory position. By adding up all products, and adding the index of the current input queue, we get the address of the exact slot.

0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3
To: 0				To: 1				To: 2				To: 3				To: 0				To: 1				To: 2				To: 3			
Phase 0																Phase 1															
VT 0																															

Figure 5.1: A single VT of the 4D array, mapped to 1D memory

5.1.4 Model and Successor Generation

Models are implemented using a data structure that represents a current state. Each state instance then provides the two operations `state_successor` and `state_violates`. The state instances are stored in the queues.

State generation is usually heavily based on branching through `switch`- and `if`-conditions. However, CUDA’s SIMT execution model, in which warps of 32 threads operate on same instruction, causes threads to pause when their execution paths diverge. This means that state generation through branching results in a slower CUDA execution. As a countermeasure, Algorithm 3 from [1] proposes to calculate all possible state transitions every time, preserving the SIMT execution and at the same time creating different states in each thread. To do so, it makes use of the custom ternary operator described in Algorithm 4 that exploits the fact that boolean values evaluate to either zero or one.

Algorithm 4 Branching-Free Ternary Operator

```
function MYTERNARY(bool c, int t, int f)
|   return  $c \cdot t + (1 - c) \cdot f$ 
```

Other than in [1], we do not fit a state’s variables into a single integer using bit shifts. Instead, we use regular C++ class data members. The increased memory demand by regular data members can be neglected, as the queues in which we store state instances are in global memory which is usually gigabytes in size. Also, the same behavior as in the paper can be achieved using C++ bit field members.

Using our custom ternary, we can re-use DVE models from the discontinued BEEM model database¹. DVE is a file format used by the DIVINE 3 model checker. We translate DVE models into CUDA compatible C++ with the template described in Algorithm 5.

5.1.5 Hash tables

The hash tables provide a data structure in which states are marked as visited. They are only used on-device within the CUDA kernel and stored in shared memory, resulting in a highly constrained size. Each VT operates on a single hash table, meaning that all threads in a block share one hash table. For memory efficiency, we use bitstate hashing, as done in [3]. Our hash tables can store $2^{18} = 262\,144$ states, taking up 32\,768 bytes of shared memory.

The hash function’s goal is to map random data with unknown distribution onto a nearly uniform distribution, so each bucket of the hash table is used equally likely. In our implementation, hashing is done using the Jenkins Hash function². The MurMurHash³, which is used by our HyperLogLog, yields similar results.

Hash collisions resulting in false-positives are intended, as explained in Section 4.3.2. To achieve that in every VT different states cause a collision, each is including a different seed into the hash. The rate of collisions can then be controlled by the hash table’s size.

In bitstate hashing, a single bit represents whether a state is visited. Per clock cycle, CUDA can transmit 32 bit of shared memory, which is stored in 32-bit words. Resulting

¹<https://paradise.fi.muni.cz/beem/>

²<https://www.burtleburtle.net/bob/hash/doobs.html>

³<https://github.com/aappleby/smhasher>

Algorithm 5 On-The-Fly State Generation on the GPU

```

▷ Global variables <
glob ← 0
▷ Process-Local variables use arrays. Here, N is the number of processes <
state[N] ← {0, ..., 0}

function state_successor(p, ndc, state)
  ▷ Evaluate all guards before calculating transitions <
  guard1 ← state[p] = 0
  guard2 ← state[p] = 1 ∧ glob ≤ 2
  guard3 ← state[p] = 1 ∧ glob > 2
  ▷ Transition from state 0 to 1 <
  next.state ← MyTernary(guard1, 1, state.state)
  ▷ Transition from state 1 to 1 <
  next.glob ← MyTernary(guard2, state.glob + 1, state.glob)
  next.state ← MyTernary(guard2, 1, next.state)
  ▷ Transition from state 1 to 0 <
  next.glob ← MyTernary(guard3, 0, next.glob)
  next.state ← MyTernary(guard3, 0, next.state)

  return next

```

from this, we use a bucket size of 32 bit. The bit representing whether a state is visited is then computed using Algorithm 6.

5.1.6 HyperLogLog++

The HyperLogLog (HLL) provides an algorithm for distributed counting of distinct elements, as described in Section 4.5. In our implementation, the grid of VTs of each run is operating on a single HLL that is stored in global memory. Visited states are added to the HLL only on-device. The host collects the HLLs after the kernels have finished, merge them into one big HLL and does the estimation of unique states visited. Merging and estimating the amount of unique states visited is done only on the host.

5.2 Usage

To verify a model using our model checker, a user has to complete the following steps:

1. Describe their model in a C++ class using Algorithm 5
2. Compile our model checker including their model
3. Execute the single output binary. On execution, the user may set the PRNG seed and number of runs using command-line options. The number of runs needed to

Algorithm 6 Bitstate hashing

```

▷ Initialize hash table buckets. Here,  $N$  is the hash table size
 $table[(1 \ll N - 5)] \leftarrow \{0, \dots, 0\}$ 
▷

function mark_visited( $state$ )
   $hash \leftarrow \text{make\_hash}(state, seed) \& (1 \ll N) - 1$ 
  ▷ First 5 bits ( $2^5 = 32$ ) are the index within the bucket
   $elem \leftarrow hash \gg (N - 5)$ 
  ▷ Last  $N - 5$  bits are the bucket index
   $bucket \leftarrow hash \& (1 \ll N - 5) - 1$ 
  ▷ Retrieve current state
   $isVisited \leftarrow table[bucket] \& (1 \ll elem) \neq 0$ 
  ▷ Update hash bucket
   $table[bucket] |= (1 \ll elem)$ 
  ▷

  return  $isVisited$ 

```

achieve a certain state space coverage needs to be evaluated experimentally.

4. Interpret the CSV output

6 Experiments

This chapter shows experimental results of our implementation. All experiments are done on an NVIDIA GeForce RTX 2080 Ti GPU, using the randomly chosen hash seed 1736331306.

The calculation of visited state space is done using a HLL with 2^{14} registers. This means that there may be a relative error of $\sigma = 1.04/\sqrt{2^{14}} = 0.008125$. The 65% error bound for each value x is then $[x \cdot (1 - \sigma), x \cdot (1 + \sigma)]$. 95% and 99% can be reached using 2σ and 3σ .

6.1 Validating our implementation and reproducing results from the paper

We start the evaluation with executing the Waypoints model introduced in Algorithm 3. Here, our model checker only implements a parallel BFS with state pruning using hash collisions. No other diversification techniques are applied.

The results of this first experiment are visualized in Figure 6.1. Both sides of the figure visualize the covered state space in relation to the executed VTs. On the left side, discovered waypoints are used as indicator. On the right side, the absolute count of the HLL divided by the total state space size is used.

The first observation that can be drawn is, that both indicators are nearly identical. Having two indicators of state space coverage, we achieve high plausibility that both the waypoints as indicators of state space coverage and the HLL as actual counting of unique visited states operate as intended.

Another observation is that even though 250 000 VTs are executed, we only cover 35% of the state space. This is far below the results from [3], where all 100 waypoints are reached after about 100 000 VTs.




Figure 6.1: Waypoints model, 1000 runs, 2^{18} hash table size, 0 start overs

There are two main implementation differences that we believe could lead to the results:

In the paper, hash tables are used that can store nearly 400 000 entries. Our hash table can only store 262 144 entries. This means that approximately $\frac{1}{3}$ fewer states can be explored per VT, resulting in less depth. The size of our hash table cannot simply be increased, even though there is shared memory left, as our bit shift based hash bucket addressing only supports table sizes that are an exponent of two.

Another reason may be that in the paper, the number of VTs per run/warp/...may be different: Initially, the paper states that they run “a single, internally parallel VT per warp”. Then, however, they state that each CUDA program “initializes a number of VTs, in this case 250”. When taking a look at the number of runs in relation to the number of discovered waypoints, results are similar to the paper again.

6.2 Introducing Start Overs

Two key observations are that the state space exploration terminates when the hash table is full and that the rate of hash table collisions and, by that, the number of states that can be visited in a VT is mainly controlled by the hash table size. This lead to the idea of *start overs*, which is based on the two-phase swarm of [3]. After a VT considers itself done,

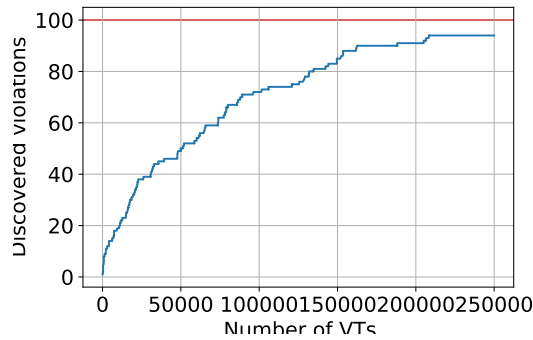


Figure 6.2: Waypoints model, 1000 runs, 2^{14} hash table size, 15 start overs

6.3 Low-Connectivity Models

The Grapple paper suspects two models of being low-connectivity: Andersons queue lock mutual exclusion algorithm and the Dining Philosophers Problem. Both models are taken from the BEEM database.

The difference in model checking performance between the *high-connectivity* Waypoints model and the *low-connectivity* Anderson and Philosophers models can be observed by comparing the number of unique states visited of the models using the same model checker configuration. Table 6.1 contains results for the three models.

Table 6.1: Comparison of three models, executed with 250 VTs (one run), 2^{14} hash table size and 10 start overs

Model	Unique States Visited	Total States Visited	Covered State Space
Anderson 3	284 941	26 688 449	-
Philosopher 15	$7.879\,01 \times 10^6$	44 674 703	55%
Waypoints	$3.183\,06 \times 10^7$	45 032 805	0.74%

6.4 Optimizing Grapple on Low-Connectivity Models

7 Conclusion

7.1 Future work

- Estimate state space size for models of unknown size - Automatically determine the optimal number of start overs

7.2 Discussion

Bibliography

- [1] Ezio Bartocci, Richard DeFrancisco, and Scott A. Smolka. “Towards a GPGPU-Parallel SPIN Model Checker”. In: *Proceedings of the 2014 International SPIN Symposium on Model Checking of Software*. SPIN 2014. San Jose, CA, USA: Association for Computing Machinery, 2014, pp. 87–96. ISBN: 9781450324526. DOI: 10.1145/2632362.2632379.
- [2] Edmund M. Clarke, Thomas A. Henzinger, and Helmut Veith. “Introduction to Model Checking”. In: *Handbook of Model Checking*. Ed. by Edmund M. Clarke et al. Cham: Springer International Publishing, 2018, pp. 1–26. ISBN: 978-3-319-10575-8. DOI: 10.1007/978-3-319-10575-8_1.
- [3] Richard DeFrancisco et al. “Swarm model checking on the GPU”. In: *International Journal on Software Tools for Technology Transfer* 22.5 (Oct. 2020), pp. 583–599. ISSN: 1433-2787. DOI: 10.1007/s10009-020-00576-x.
- [4] Philippe Flajolet et al. “HyperLogLog: the analysis of a near-optimal cardinality estimation algorithm”. In: *Discrete Mathematics & Theoretical Computer Science DMTCS Proceedings* vol. AH, 2007 Conference on Analysis of Algorithms (AofA 07) (Jan. 2007). DOI: 10.46298/dmtcs.3545.
- [5] Stefan Heule, Marc Nunkesser, and Alexander Hall. “HyperLogLog in Practice: Algorithmic Engineering of a State of the Art Cardinality Estimation Algorithm”. In: *Proceedings of the 16th International Conference on Extending Database Technology*. EDBT ’13. Genoa, Italy: Association for Computing Machinery, 2013, pp. 683–692. ISBN: 9781450315975. DOI: 10.1145/2452376.2452456.
- [6] Gerard J. Holzmann. “Explicit-State Model Checking”. In: *Handbook of Model Checking*. Ed. by Edmund M. Clarke et al. Cham: Springer International Publishing, 2018, pp. 153–171. ISBN: 978-3-319-10575-8. DOI: 10.1007/978-3-319-10575-8_5.
- [7] Gerard J. Holzmann. “Parallelizing the Spin Model Checker”. In: *Model Checking Software*. Ed. by Alastair Donaldson and David Parker. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 155–171. ISBN: 978-3-642-31759-0. DOI: 10.1007/978-3-642-31759-0_12.
- [8] Gerard J. Holzmann, Rajeev Joshi, and Alex Groce. “Swarm Verification”. In: *2008 23rd IEEE/ACM International Conference on Automated Software Engineering*. 2008 23rd IEEE/ACM International Conference on Automated Software Engineering. 2008, pp. 1–6. DOI: 10.1109/ASE.2008.9.
- [9] Gerard J. Holzmann, Rajeev Joshi, and Alex Groce. “Swarm Verification Techniques”. In: *IEEE Transactions on Software Engineering* 37.6 (2011), pp. 845–857. DOI: 10.1109/TSE.2010.110.